

# CS598 PSL — Assignment 1 — Problem 1

Group Members: Rahul Kasibhatla, Neeyati Devanagondi, Manas Gandhi

## Question 1

```
In [1]: #Question 1a
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

def gap_statistic(X, refs=10, max_clusters=10, random_state=42):
    np.random.seed(random_state)
    shape = X.shape
    tops = X.max(axis=0)
    bottoms = X.min(axis=0)
    dists = np.zeros((refs, shape[0], shape[1]))

    for i in range(refs):
        dists[i] = np.random.uniform(bottoms, tops, size=shape)

    gaps = []
    for k in range(1, max_clusters + 1):
        km = KMeans(n_clusters=k, random_state=random_state).fit(X)
        disp = km.inertia_

        ref_disps = np.zeros(refs)
        for i in range(refs):
            km_ref = KMeans(n_clusters=k, random_state=random_state).fit(dists[i])
            ref_disps[i] = km_ref.inertia_

        gap = np.log(np.mean(ref_disps)) - np.log(disp)
        gaps.append(gap)

    optimal_k = np.argmax(gaps) + 1
    return optimal_k, gaps

df = pd.read_csv("/content/Mall_Customers.csv")

df_encoded = pd.get_dummies(df, columns=["Genre"], drop_first=True)

features = ["Age", "Annual Income (k$)", "Spending Score (1-100)", "Genre_Ma
X = df_encoded[features]
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

n_clusters_gap, gaps = gap_statistic(X_scaled, refs=10, max_clusters=10)
print("Optimal k (Gap Statistic):", n_clusters_gap)

plt.plot(range(1, 11), gaps, marker='o')
plt.xlabel("k")
plt.ylabel("Gap Value")
plt.title("Gap Statistic")
plt.show()

sil_scores = []
K = range(2, 11)

for k in K:
    km = KMeans(n_clusters=k, random_state=42)
    labels = km.fit_predict(X_scaled)
    sil_scores.append(silhouette_score(X_scaled, labels))

ksil = K[np.argmax(sil_scores)]
print("Optimal k (Silhouette Score):", ksil)

plt.plot(K, sil_scores, marker="o")
plt.xlabel("k")
plt.ylabel("Silhouette Score")
plt.title("Silhouette Scores")
plt.show()

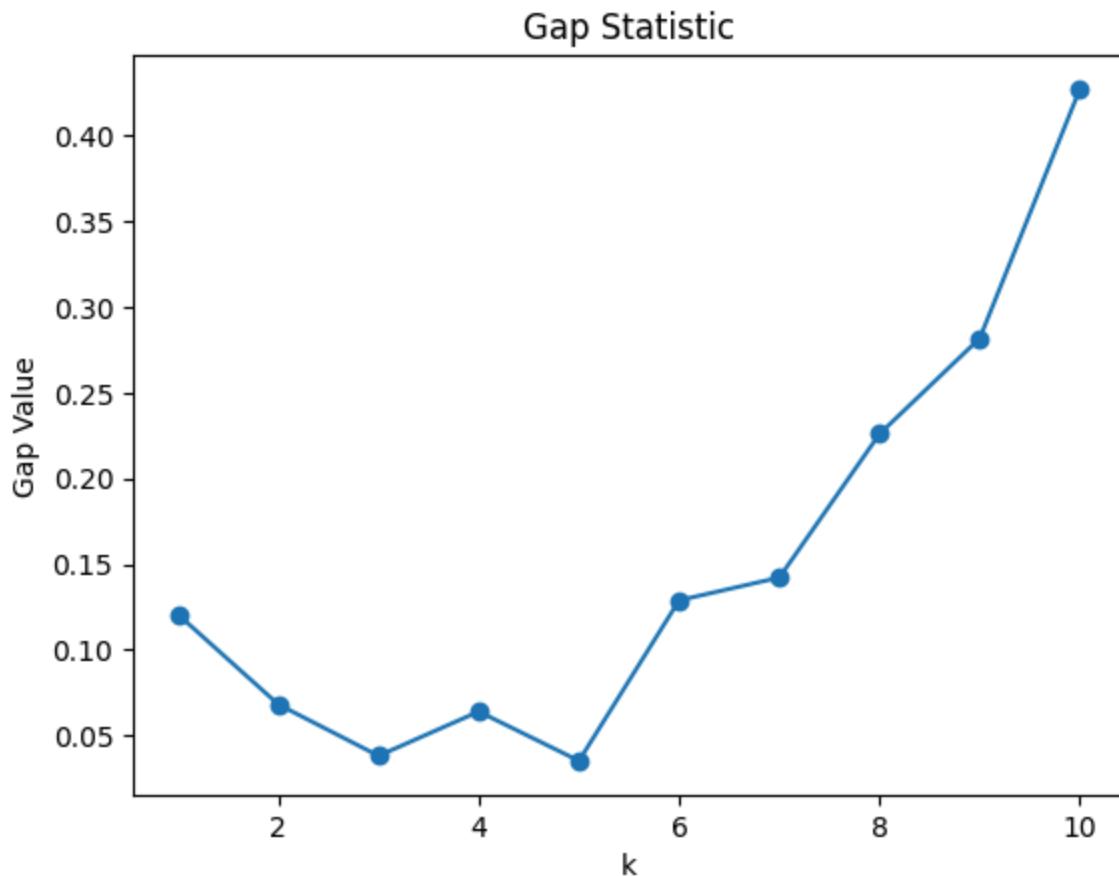
k_values = [n_clusters_gap, ksil]

for k in k_values:
    km = KMeans(n_clusters=k, random_state=42)
    df_encoded[f"Cluster_k{k}"] = km.fit_predict(X_scaled)

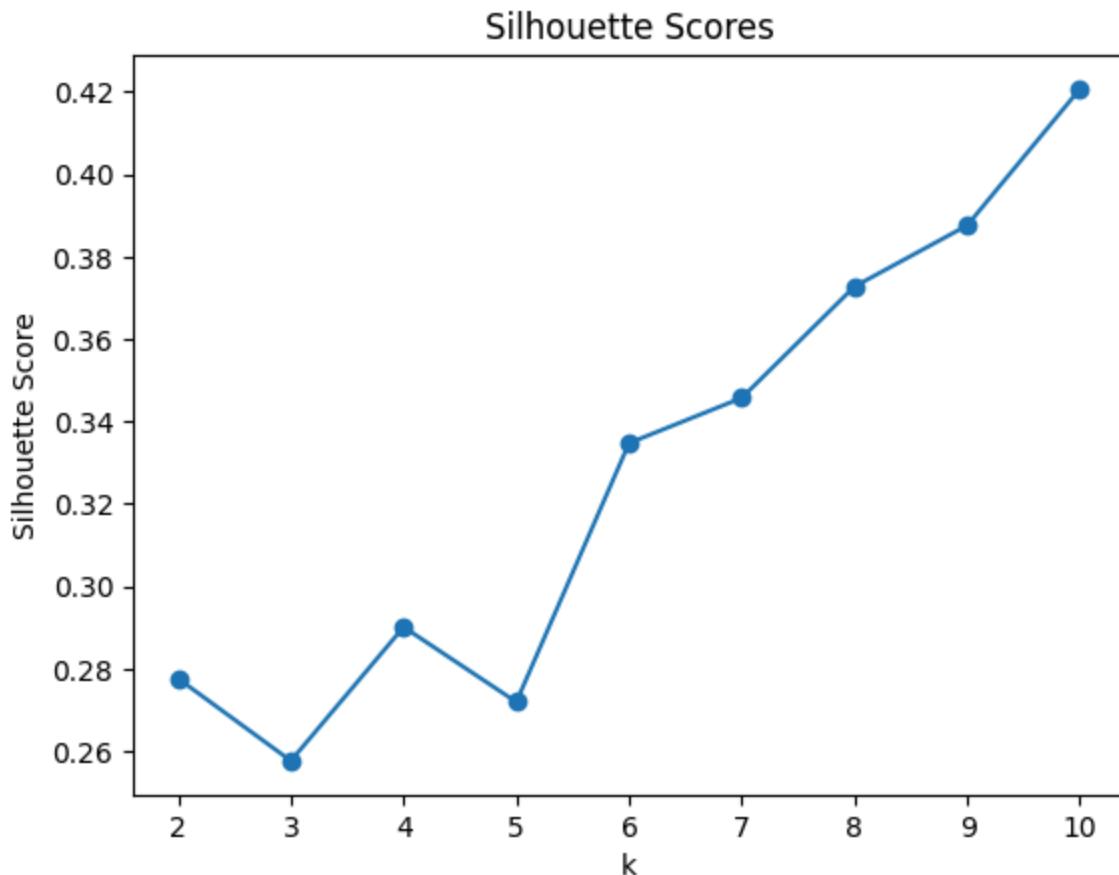
for k in k_values:
    sns.scatterplot(
        data=df_encoded,
        x="Annual Income (k$)",
        y="Spending Score (1-100)",
        hue=f"Cluster_k{k}"
    )
    plt.title(f"K-Means Clusters (k={k})")
    plt.show()

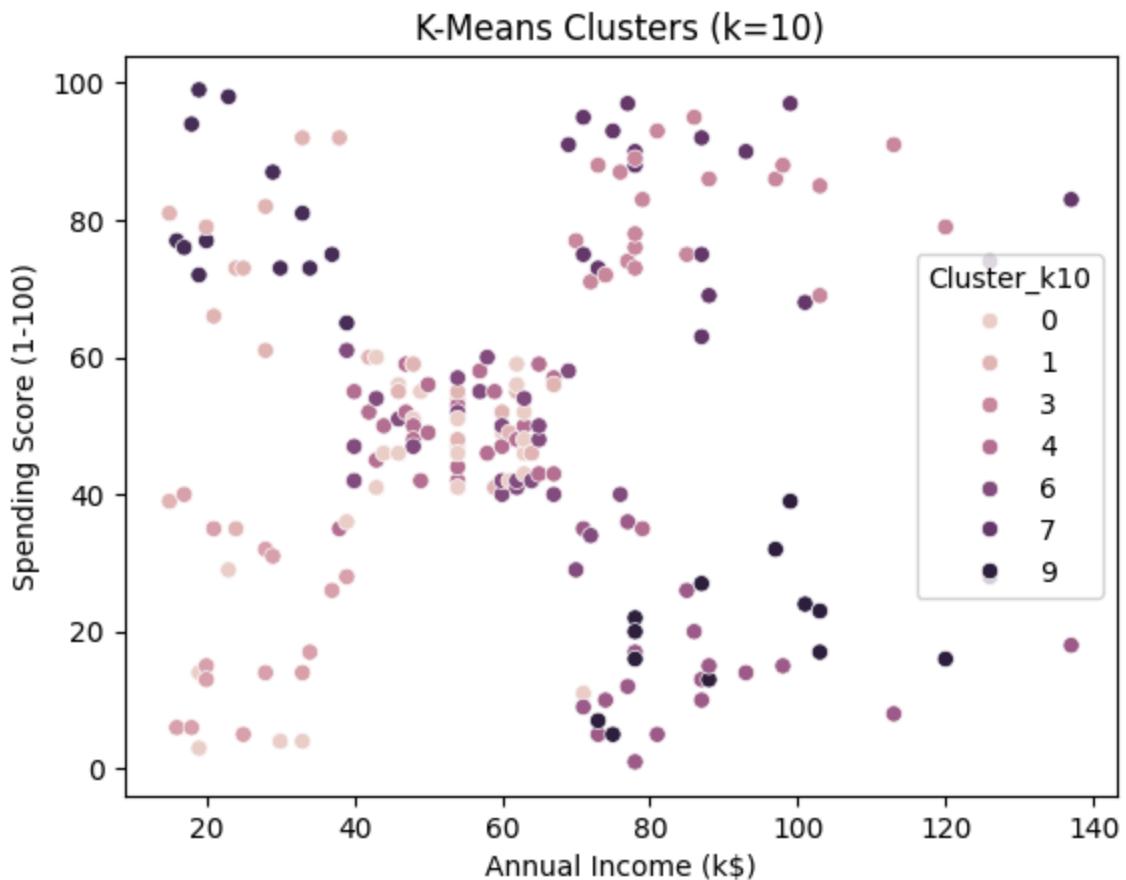
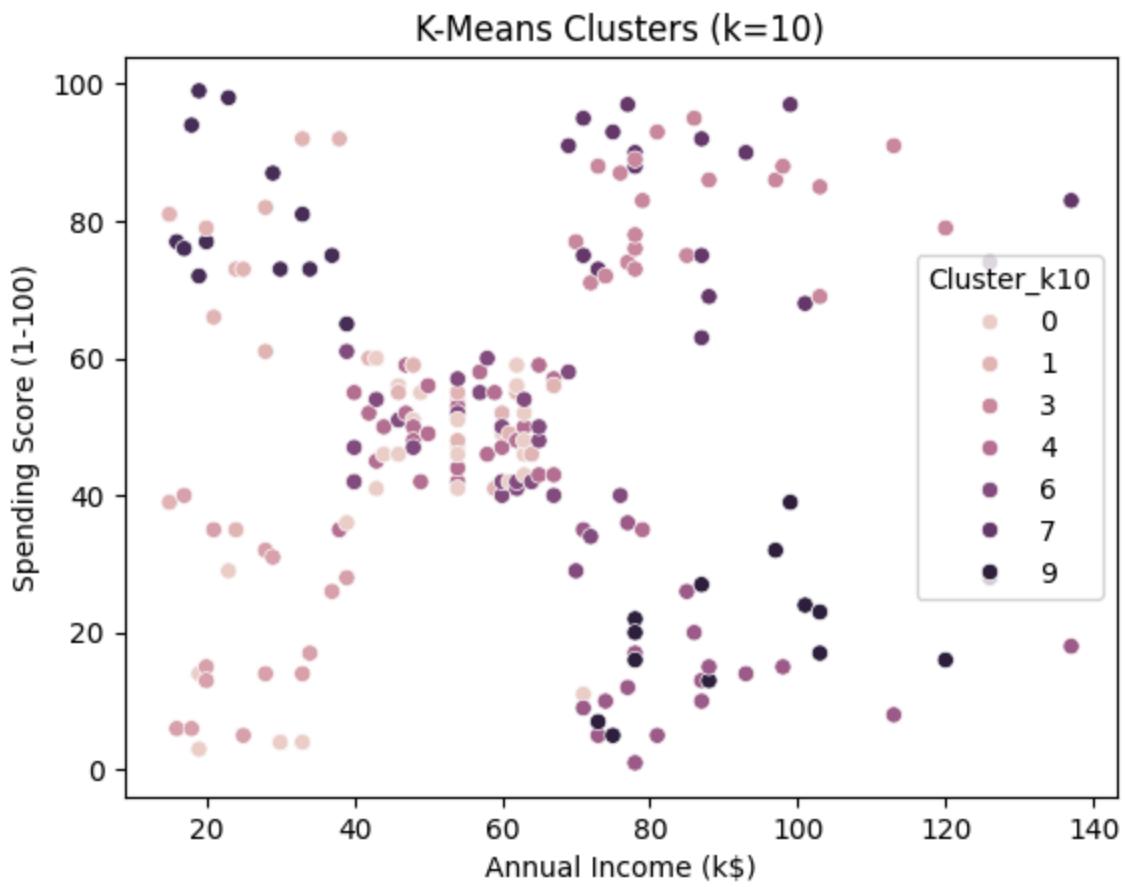
for k in k_values:
    print("\n Summary for k =", k, "")
    print(
        df_encoded.groupby(f"Cluster_k{k}")[
            ["Age", "Annual Income (k$)", "Spending Score (1-100)", "Genre_M"]
        ].mean()
    )
```

Optimal k (Gap Statistic): 10



Optimal k (Silhouette Score): 10





Summary for k = 10

	Age	Annual Income (k\$)	Spending Score (1-100)	Genre_Ma le
<b>Cluster_k10</b>				
0	58.846154	48.692308	39.846154	1.0000
00	25.250000	41.250000	60.916667	1.0000
1	41.214286	26.071429	20.142857	0.0714
29	32.190476	86.047619	81.666667	0.0000
3	54.153846	54.230769	48.961538	0.0000
00	38.473684	85.894737	14.210526	1.0000
5	27.960000	57.360000	47.120000	0.0000
00	33.277778	87.111111	82.666667	1.0000
6	25.461538	25.692308	80.538462	0.0000
00	43.785714	93.285714	20.642857	0.0000
9				
00				

Summary for k = 10

	Age	Annual Income (k\$)	Spending Score (1-100)	Genre_Ma le
<b>Cluster_k10</b>				
0	58.846154	48.692308	39.846154	1.0000
00	25.250000	41.250000	60.916667	1.0000
1	41.214286	26.071429	20.142857	0.0714
29	32.190476	86.047619	81.666667	0.0000
3	54.153846	54.230769	48.961538	0.0000
00	38.473684	85.894737	14.210526	1.0000
5	27.960000	57.360000	47.120000	0.0000
00	33.277778	87.111111	82.666667	1.0000
6	25.461538	25.692308	80.538462	0.0000
00	43.785714	93.285714	20.642857	0.0000
9				
00				

1a.

ii

The scatterplots for k = 10 clear regions where spending score and income combine to form customer groups. High-income/high-spending customers cluster cleanly in the upper right region, while low-income/low-spending customers cluster in the lower left. However, as k increases, some clusters become more granular and begin to overlap, suggesting that while k = 10 does separate major behavior patterns, it also creates smaller subclusters that are not perfectly separable in a 2d space. Overall, the clusters show reasonable separability but with some overlap between middle-income, moderate-spending groups.

### iii

Overall, the k = 10 clustering reveals that customers are primarily differentiated by their income and spending patterns. Younger customers tend to show more impulsive spending behavior, while older customers generally spend less and follow more conservative patterns. There are also noticeable gender differences within some clusters, particularly among the highest and lowest spenders.

```
In [13]: #Question 1b

from scipy.cluster.hierarchy import dendrogram, linkage, fcluster

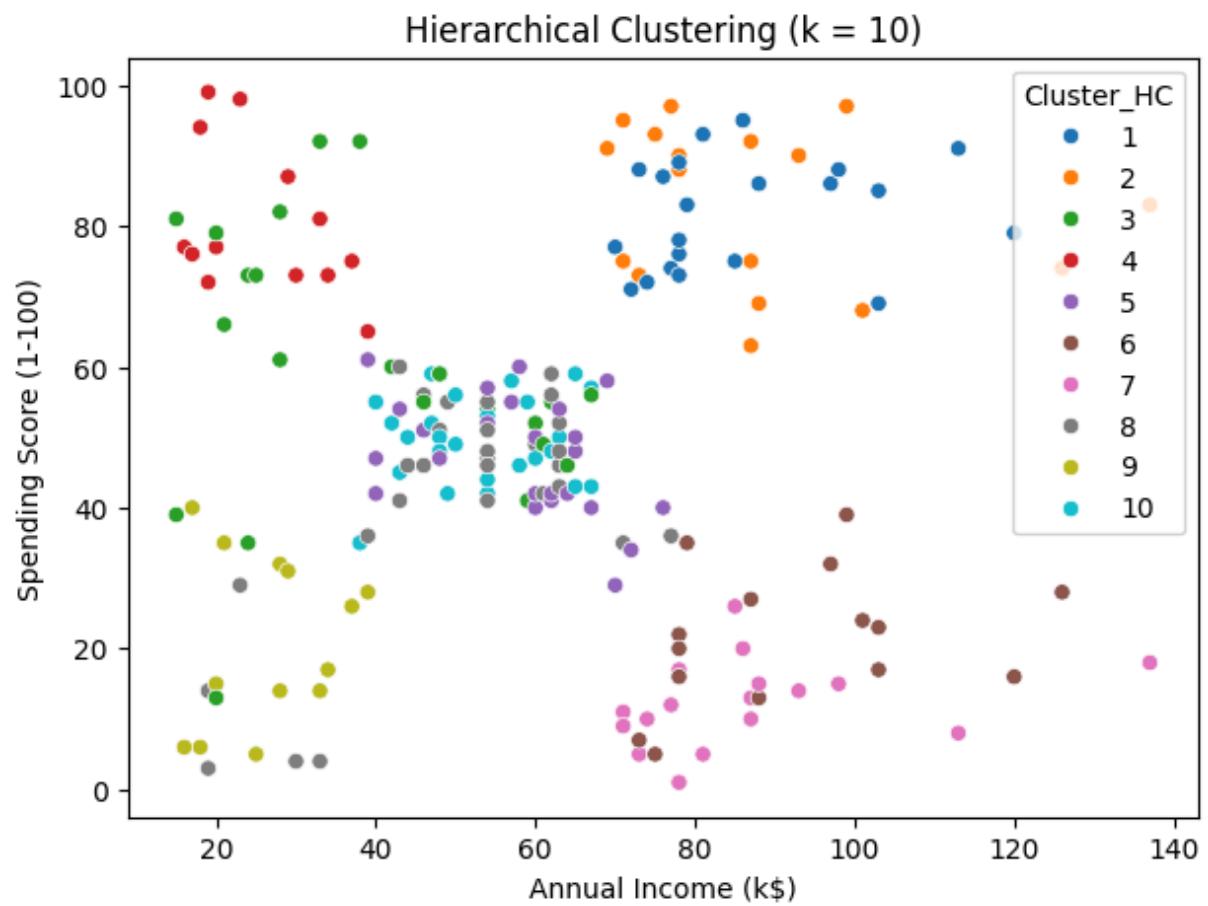
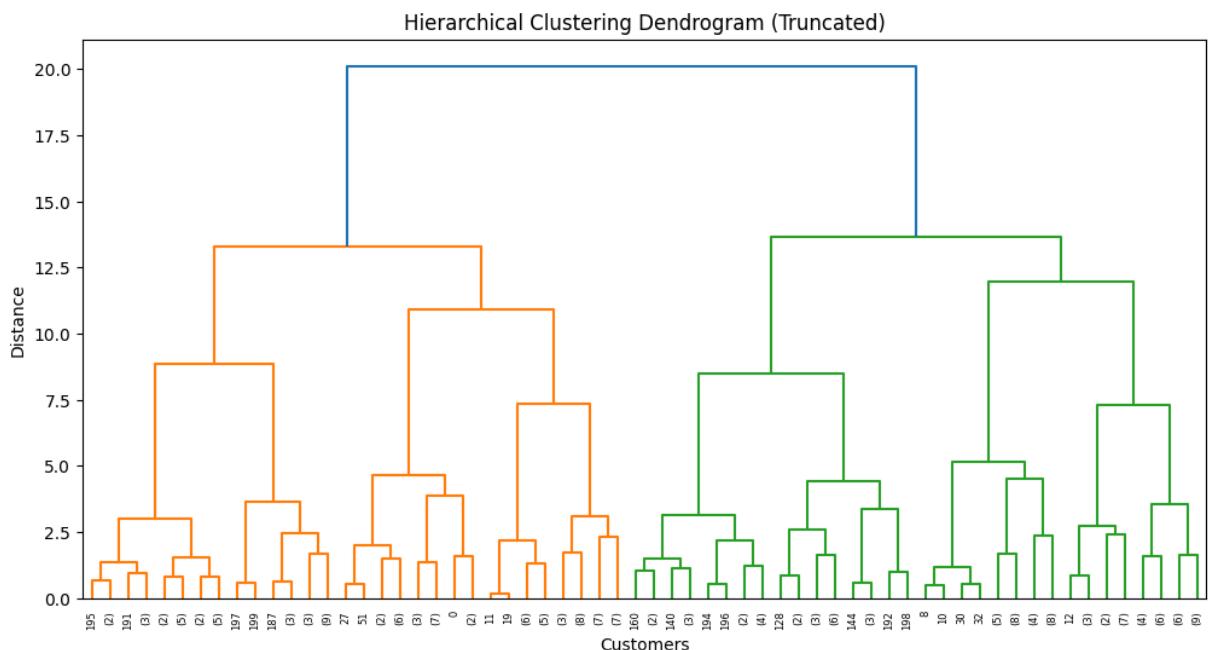
Z = linkage(X_scaled, method='ward')

plt.figure(figsize=(12, 6))
dendrogram(Z, truncate_mode='level', p=5)
plt.title("Hierarchical Clustering Dendrogram (Truncated)")
plt.xlabel("Customers")
plt.ylabel("Distance")
plt.show()

k_hier = 10
df_encoded["Cluster_HC"] = fcluster(Z, k_hier, criterion='maxclust')

plt.figure(figsize=(7,5))
sns.scatterplot(
    data=df_encoded,
    x="Annual Income (k$)",
    y="Spending Score (1-100)",
    hue="Cluster_HC",
    palette="tab10"
)
plt.title("Hierarchical Clustering (k = 10)")
plt.show()

print(df_encoded.groupby("Cluster_HC")[
    ["Age", "Annual Income (k$)", "Spending Score (1-100)", "Genre_Male"]
].mean())
```



	Age	Annual Income (k\$)	Spending Score (1-100)	Genre_Male
e	Cluster_HC			
1	32.190476	86.047619	81.666667	0.
0	33.277778	87.111111	82.666667	1.
2	24.565217	39.217391	59.652174	1.
0	25.461538	25.692308	80.538462	0.
4	27.960000	57.360000	47.120000	0.
0	44.600000	92.333333	21.600000	0.
6	38.833333	86.388889	11.666667	1.
0	56.551724	50.034483	41.344828	1.
8	41.538462	26.538462	20.692308	0.
0	54.080000	53.240000	49.520000	0.
10				
0				

## 1b.

ii

The hierarchical clustering scatterplot shows similar overall patterns to K-Means, with clear high-income/high-spending and low-income/low-spending groups. However, the hierarchical clusters are less compact and show more overlap because hierarchical clustering does not optimize cluster tightness the way K-Means does. K-Means produces cleaner, more separated clusters, while hierarchical clustering yields more irregular shapes and cluster sizes.

iii

The hierarchical clustering groups mirror the behavioral patterns seen in K-Means. There are high-income/high-spending clusters (Clusters 1–2), low-income/high-spending shoppers (Cluster 4), and high-income but low-spending customers (Clusters 6–7). Middle-income moderate spenders also form several clusters (5, 8, 10). Overall, hierarchical clustering identifies similar customer segments but organizes them slightly differently due to the linkage-based clustering process.

## Question 2

In [14]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.cluster import KMeans, SpectralClustering

np.random.seed(598)

n = 200
n_per_moon = n // 2
r = 2
noise_var = 0.05
noise_std = np.sqrt(noise_var)

theta1 = np.random.uniform(0, np.pi, n_per_moon)
x1 = r * np.cos(theta1) + np.random.normal(0, noise_std, n_per_moon)
y1 = r * np.sin(theta1) + np.random.normal(0, noise_std, n_per_moon)
label1 = np.zeros(n_per_moon, dtype=int)

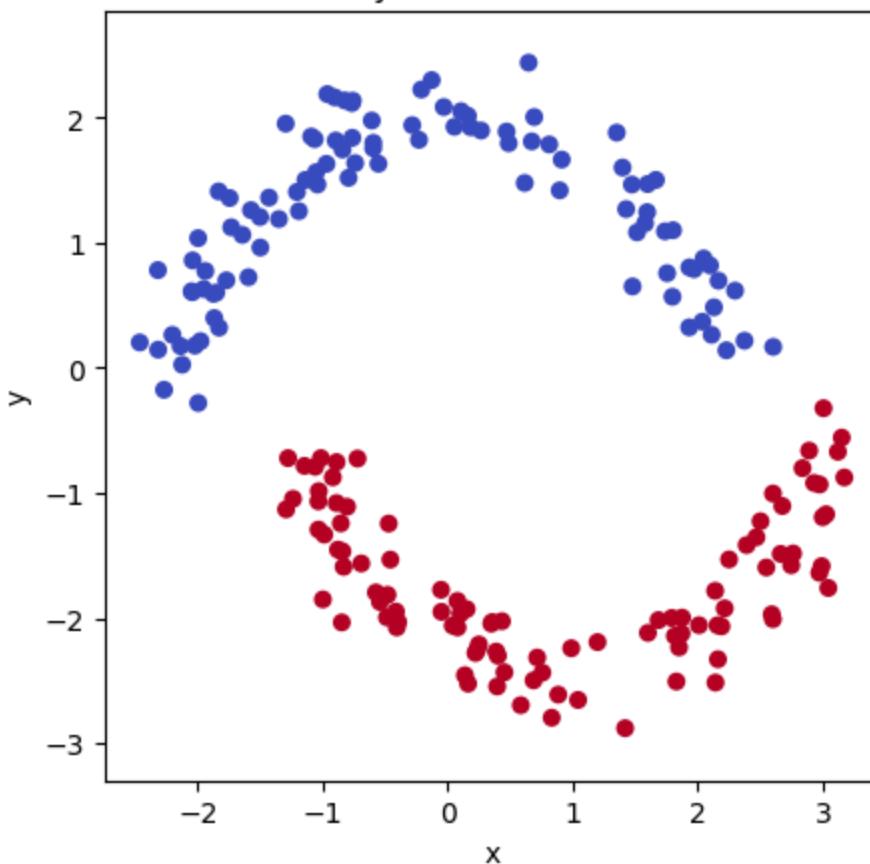
theta2 = np.random.uniform(0, np.pi, n_per_moon)
x2 = 1 - r * np.cos(theta2) + np.random.normal(0, noise_std, n_per_moon)
y2 = -r * np.sin(theta2) - 0.5 + np.random.normal(0, noise_std, n_per_moon)
label2 = np.ones(n_per_moon, dtype=int)

X = np.vstack([np.column_stack([x1, y1]), np.column_stack([x2, y2])])
y = np.concatenate([label1, label2])

df_moons = pd.DataFrame({"x": X[:, 0], "y": X[:, 1], "label": y})

plt.figure(figsize=(5,5))
plt.scatter(df_moons["x"], df_moons["y"], c=df_moons["label"], cmap="coolwarm")
plt.title("Two Moons Synthetic Data (True Labels)")
plt.xlabel("x")
plt.ylabel("y")
plt.axis("equal")
plt.show()
```

### Two Moons Synthetic Data (True Labels)



```
In [15]: # 2a
kmeans = KMeans(n_clusters=2, random_state=42, n_init=10)
df_moons["kmeans_cluster"] = kmeans.fit_predict(X)

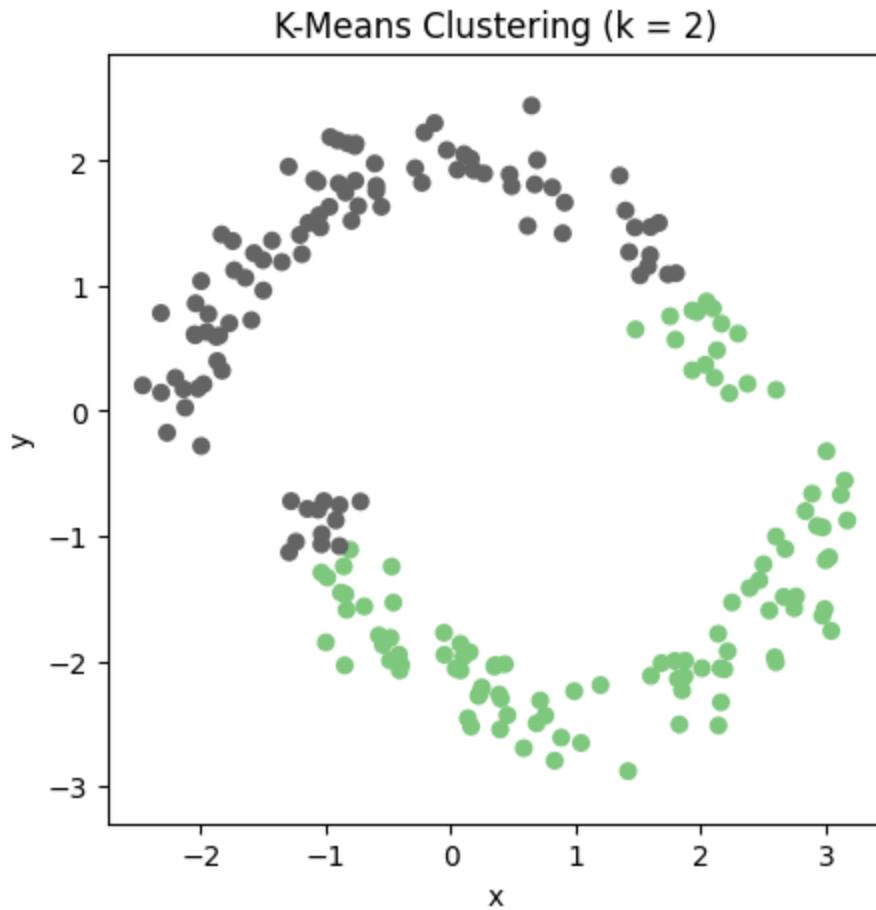
plt.figure(figsize=(5,5))
plt.scatter(df_moons["x"], df_moons["y"],
            c=df_moons["kmeans_cluster"], cmap="Accent", s=30)
plt.title("K-Means Clustering (k = 2)")
plt.xlabel("x")
plt.ylabel("y")
plt.axis("equal")
plt.show()

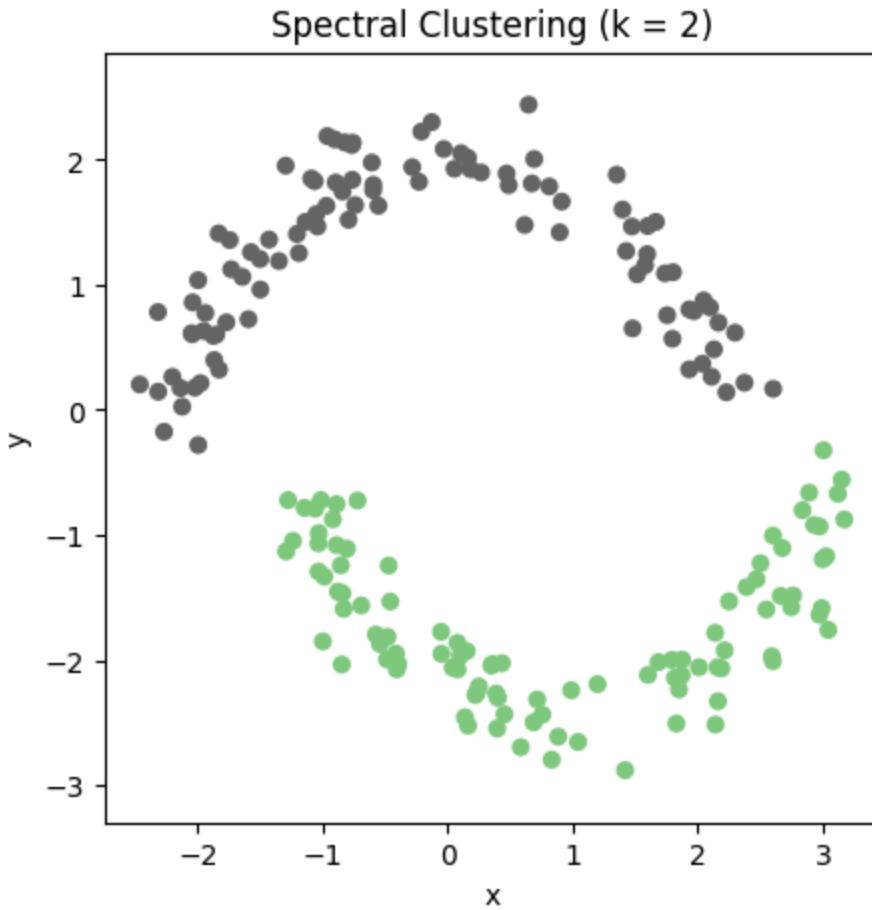
#2b

spec = SpectralClustering(
    n_clusters=2,
    affinity="rbf",
    assign_labels="kmeans",
    random_state=42
)
df_moons["spectral_cluster"] = spec.fit_predict(X)

plt.figure(figsize=(5,5))
plt.scatter(df_moons["x"], df_moons["y"],
            c=df_moons["spectral_cluster"], cmap="Accent", s=30)
plt.title("Spectral Clustering (k = 2)")
```

```
plt.xlabel("x")
plt.ylabel("y")
plt.axis("equal")
plt.show()
```





## 2a

Using K-Means with k=2 on the two-moons data does not correctly recover the true structure. Because K-Means assumes clusters are roughly spherical, it places a straight decision boundary through the middle of the dataset. As a result, the two curved moon shapes are mixed: some points from the upper moon are grouped with the lower moon and vice versa. The algorithm fails to capture the non-linear geometry of the moons, leading to poor separation between the two true groups.

## 2b

Spectral clustering with an RBF kernel has two moon shapes. The resulting clusters align almost perfectly with the true labels, cleanly separating the top and bottom moon structures. Compared to part (a), spectral clustering handles the non-linear boundary correctly and produces a more accurate clustering.